

UiT

THE ARCTIC
UNIVERSITY
OF NORWAY

Protected Mode

A brief introduction

Erlend Graff

INF-2201 Operating System Fundamentals

Department of Computer Science

University of Tromsø



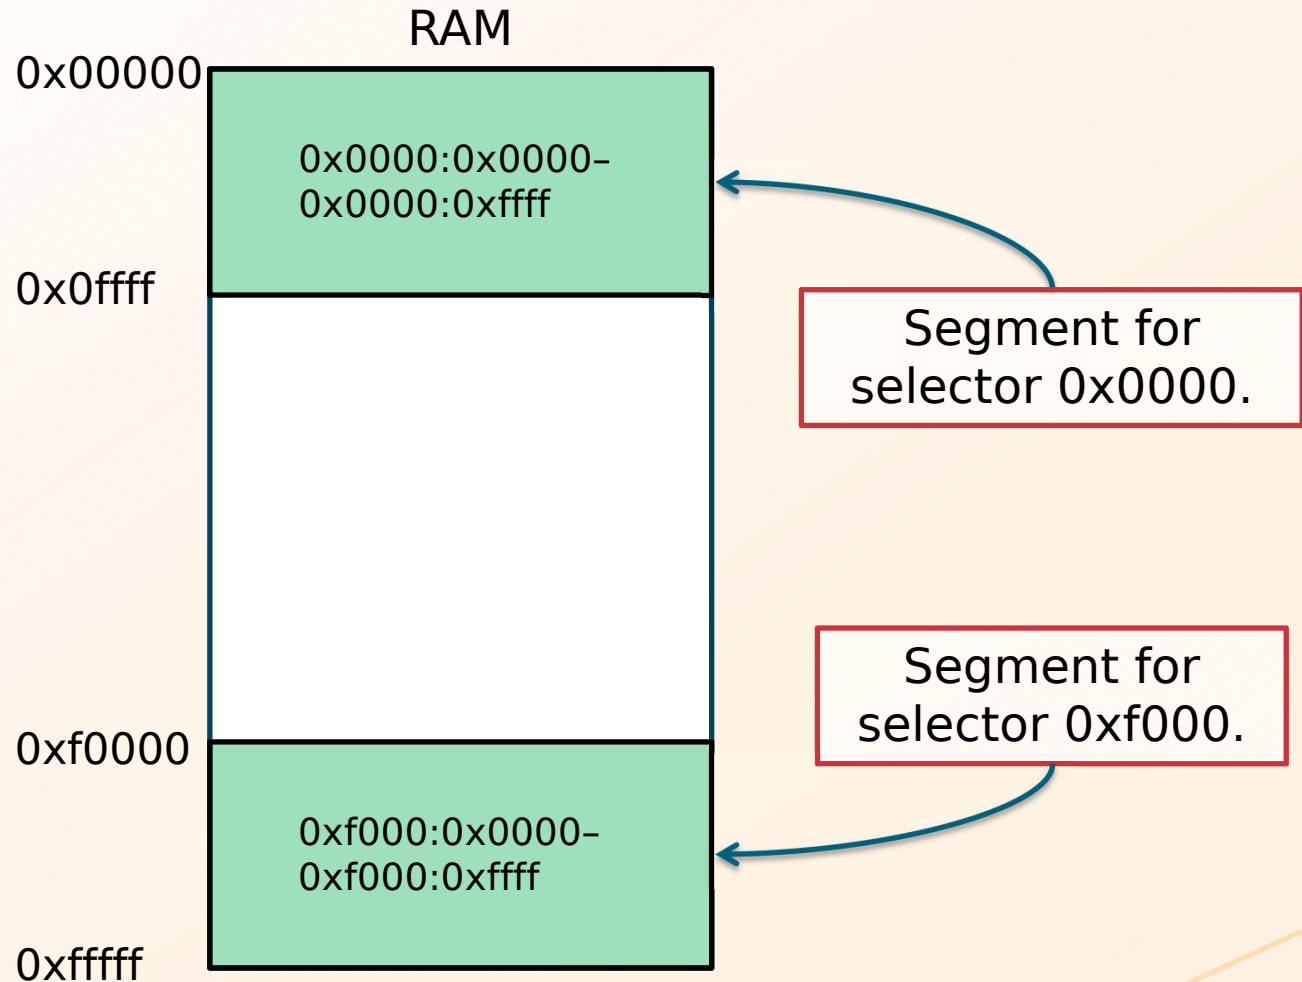
Real Mode - recap

- Legacy mode - Intel 8086 architecture backwards-compatibility
- 16-bit segmented memory model
- 1MB memory (20-bit linear address space)

Real Mode segmentation

- 1MB memory space, 20-bit linear address space.
- `0x00000–0xffffffff` should be addressable range.
- 16-bit logical addresses.
 - Each segment is 64KB.
- Address translation
 - $\text{Linear address} = (\text{Selector} \ll 4) + \text{Offset}$
- Examples
 - `0x0000:0x0000 = 0x00000`
 - `0x0000:0xffff = 0x0ffff`
 - `0xf000:0x0000 = 0xf0000`
 - `0xf000:0xffff = 0xfffffff`

Real Mode segmentation - example



Uh...

But what should we do about it?

A problem with segmentation (1)

- Selector 0x0000 addresses range 0x00000–0x0ffff.
- Selector 0xf000 addresses range 0xf0000–0xfffff.
- What about selectors 0xf001–0xffff?
- Shouldn't 0xffff:0xffff be equal to 0x10ffef?
 - But 0x10ffef is outside of 1MB range 0x00000–0xfffff.

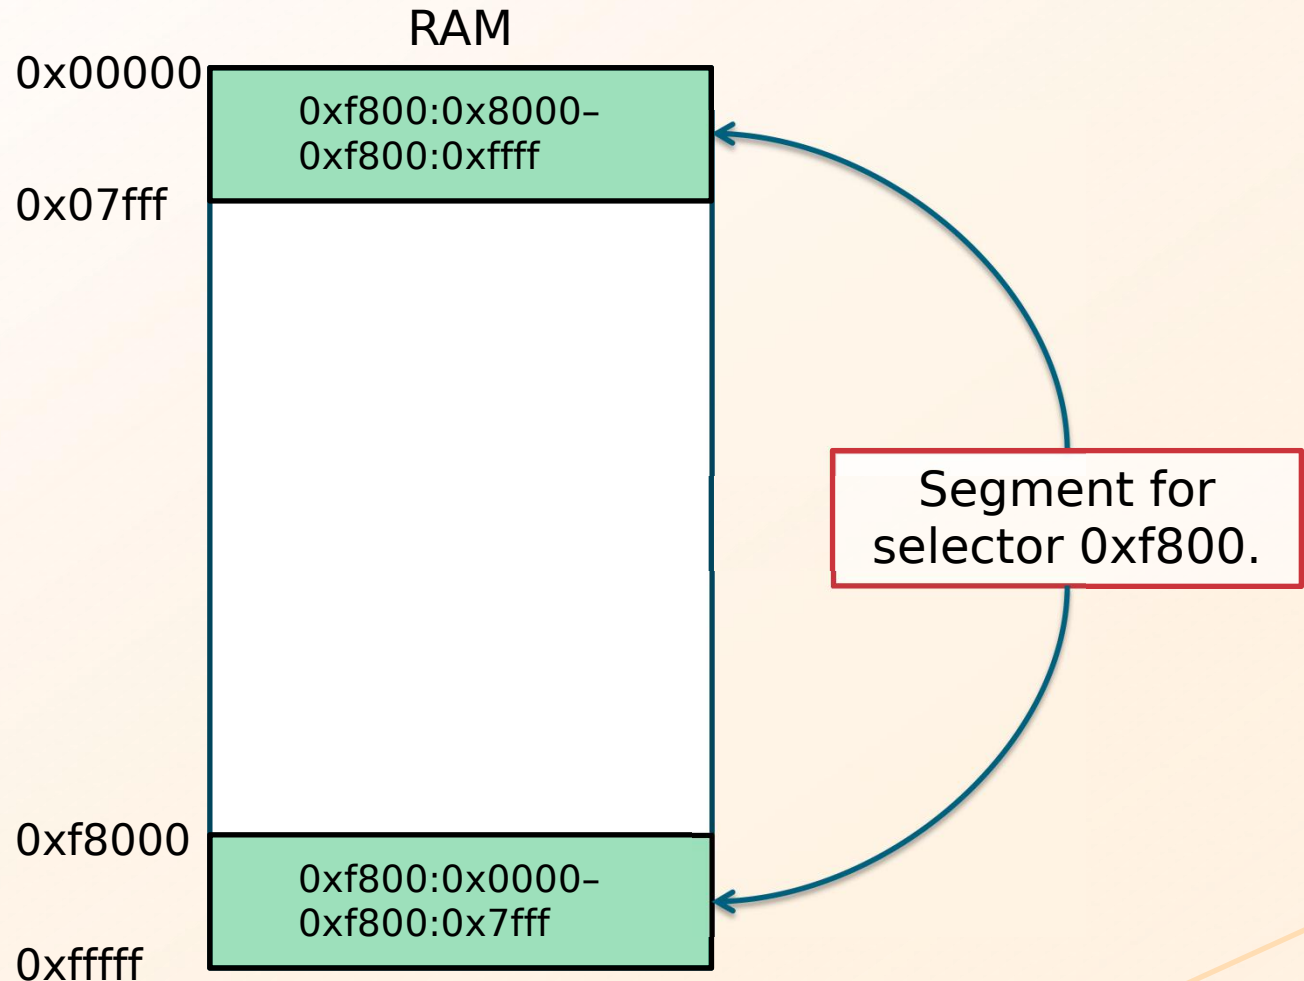
A problem with segmentation (2)

- 8086 had 20 address lines: A0–A19.
 - This is the reason for having 1MB (20-bit) linear address space in Real Mode.
 - Would be impossible to address memory outside 1MB range 0x00000–0xfffff.
 - Segments that cross 0xffff wrap around. In essence, normal calculation of linear address, but with overflowing 20-bit addition.
 - E.g. 0xffff:0xffff = 0x10ffef = 0x0ffef.




Overflow

Segmentation - example of wraparound




A problem with segmentation (3)

- 8086 had 20 address lines: A0–A19.
 - This is the reason for having 1MB (20-bit) linear address space in Real Mode.
 - Would be impossible to address memory outside 1MB range 0x00000–0xfffff.
 - Segments that cross 0xffff wrap around. In essence, normal calculation of linear address, but with overflowing 20-bit addition.
 - E.g. 0xffff:0xffff = 0x10ffef = 0x0ffef.
- But 80286 had 24 address lines: A0–A23.
 - Could address 16MB (24-bit linear address space).
 - Still 64KB segments (16-bit logical addresses).
 - But segments crossing 1MB boundary do not wrap around!



A problem with segmentation (4)

- Intel 80286 would break 8086 Real Mode backwards compatibility.
- To prevent this, the 80486 introduced the A20M pin.
 - The A20M pin is asserted by default to disable the A20 address line and enforce wraparound.
 - It is part of the 8042 keyboard controller, because it (unfortunately) had a spare pin.



A problem with segmentation (4)

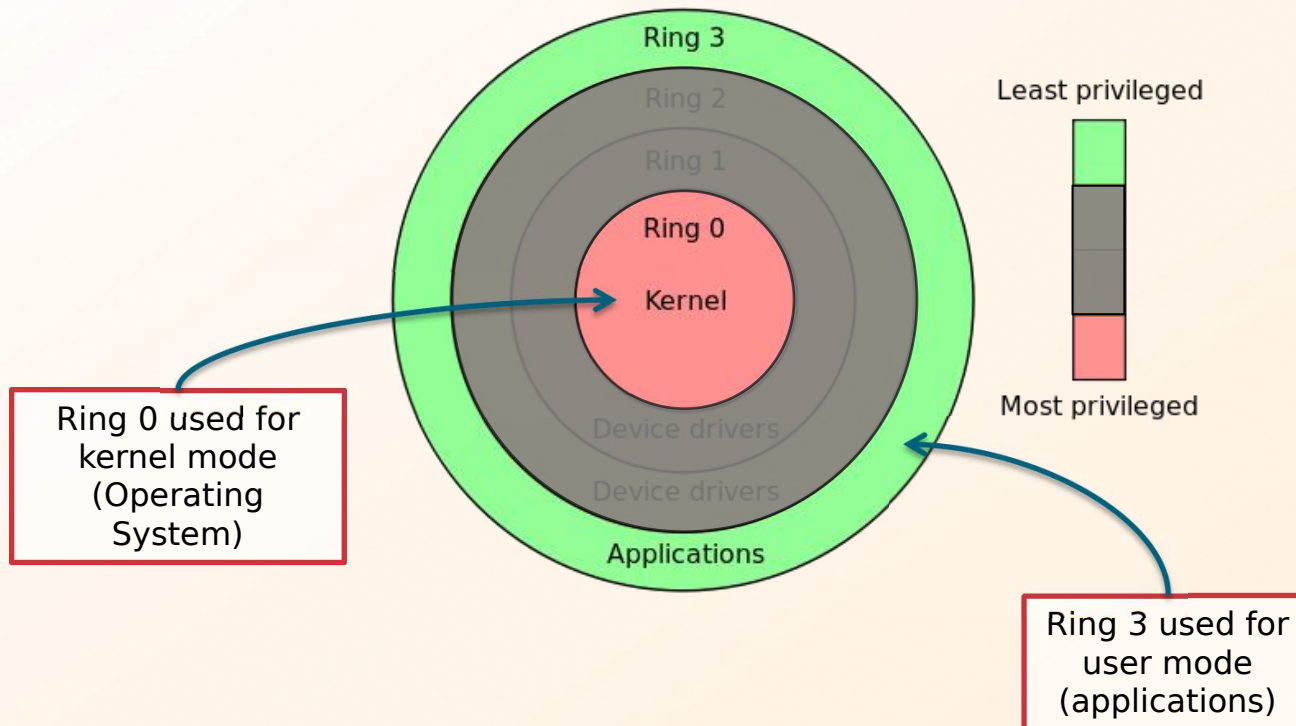
- Intel 80286 would break 8086 Real Mode backwards compatibility.
- To prevent this, the 80486 introduced the A20M pin.
 - The A20M pin is asserted by default to disable the A20 address line and enforce wraparound.
 - It is part of the 8042 keyboard controller, because it (unfortunately) had a spare pin.
- The pin must be asserted while in Real Mode, but may be deasserted to enable the A20 address line before changing processor mode.

Protected Mode

- Problems with Real Mode are the “features” of Real Mode
 - 1MB memory – too little
 - Segmented memory model – tedious
 - 16-bit is sooo 1978!
 - No protection
 - No isolation
 - No virtual memory
 - Etc.
- Protected mode
 - Introduced with 80286 as 16-bit operating mode.
 - Today’s 32-bit Protected Mode came with 80386.
 - Protected mode is the most widely used operating mode of 32-bit x86 processors.
 - Some features:
 - 4GB memory space.
 - Virtual memory / paging mechanisms with memory access rights.
 - 4 separate privilege levels for executing instructions.
 - Mechanisms for safe trapping and multi-tasking.

Protected Mode privilege levels

- Protected Mode offers 4 privilege levels
- Only 2 are used in practice!



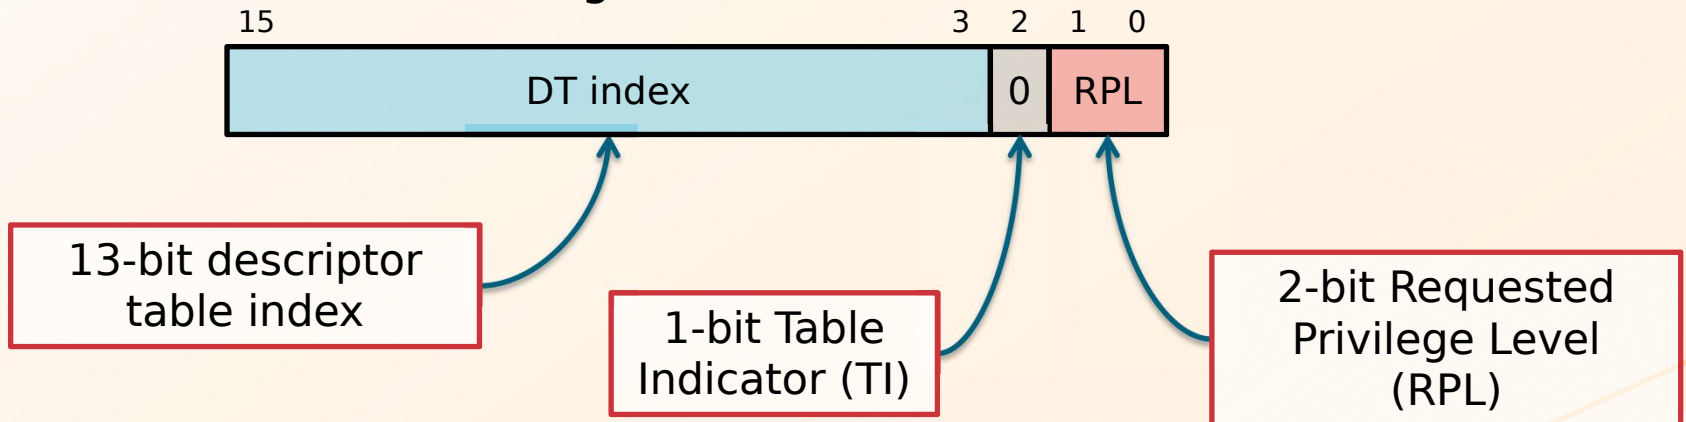
Protected Mode segmentation (1)

- More complex than in Real Mode.
- Still allows for a segmented memory model, but this is not used in practice. Instead, segments are set equal to entire 4GB address space.
- Then, why do we need segments?
 - Protection / access rights!
- Segments are no longer of fixed size. Instead, have both base and limit (size of segment).

Protected Mode segmentation (2)

- Segment registers are still in use
 - But segment registers no longer store segment base addresses in encoded form.
 - Segments registers consist of:
 - 16-bit selector (“visible”, may be written to and read from segment register).
 - 32-bit segment base address (“hidden”, in descriptor cache).
 - 32-bit segment limit (“hidden”, in descriptor cache).
 - 32-bit access rights (“hidden”, in descriptor cache).

16-bit segment selector



Protected Mode segmentation (3)

- Segment selector
 - The Requested Privilege Level (RPL) specifies the most privileged level (ring) that is allowed when loading a new segment (the segment that is loaded may only be equally or less privileged than the RPL). It corresponds to the ring number (0-3), so lower is more privileged.
 - The CS segment selector holds the Current Privilege Level (CPL) instead of an RPL.
 - The 13 MSB specifies the index into a descriptor table where a descriptor for the segment is found.
 - The Table Indicator specifies in which of two possible descriptor tables the segment descriptor is located:
 - TI = 0 - the descriptor is found in the Global Descriptor Table (GDT)
 - TI = 1 - the descriptor is found in the Local Descriptor Table (LDT)
 - In practice, the Local Descriptor Table is rarely used.

Global Descriptor Table (1)

- A fundamental data structure in x86 architecture.
- Only one instance of GDT!
- Contains 8-byte descriptors
 - One null descriptor (first descriptor, required but unused)
 - Segment descriptors
 - At least two – one for code and one for data.
 - Typically four – one pair of code and data descriptors for each privilege level.
 - Other descriptor types
 - Task State Segment (TSS) descriptors – at least (and most typically) one per CPU.
 - Local Descriptor Table (LDT) descriptors (rarely used anymore).
 - Call Gate descriptors (rarely used anymore).

Global Descriptor Table (2)



- Segment descriptors
 - Consist of
 - Segment base address (32 bit) – should be 0x00000000.
 - Segment limit (20 bit) – should be 0xffff.
 - Flags
 - Access byte



Global Descriptor Table (3)

- Flags

- Granularity bit (G)
 - G = 0 - limit is in bytes, G = 1 - limit is in 4KB blocks. Should be 1.
- Default Size bit (D/B) - should be 1 for 32-bit protected mode (0 used for 16-bit PM).
- Available for System bit (A) - should always be 0.



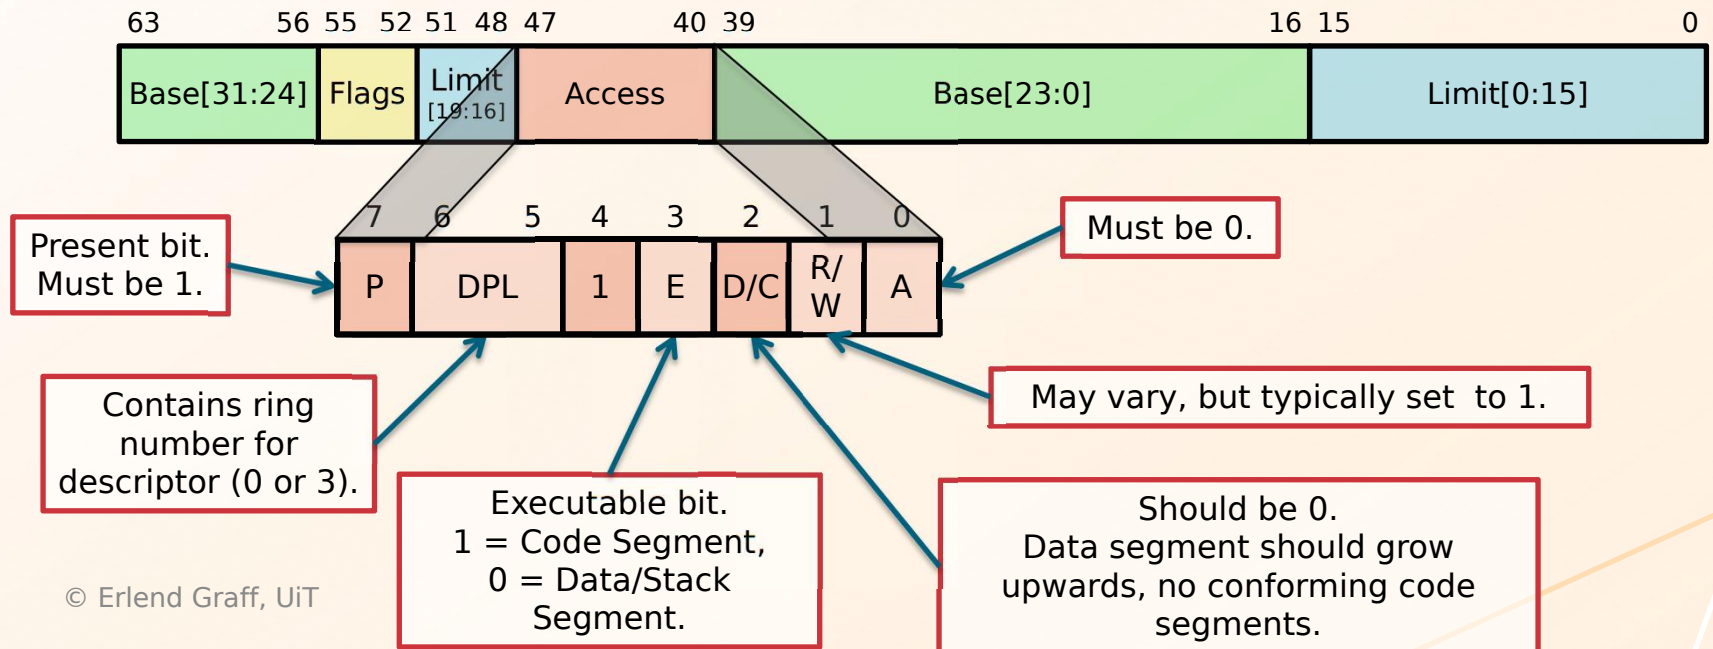
Granularity bit.
G = 1, Base = 0x0 and
Limit = 0xffff will make
segment span 4GB.

Default Size bit.
1 = 32-bit PM.

Available for System bit.
Always 0.

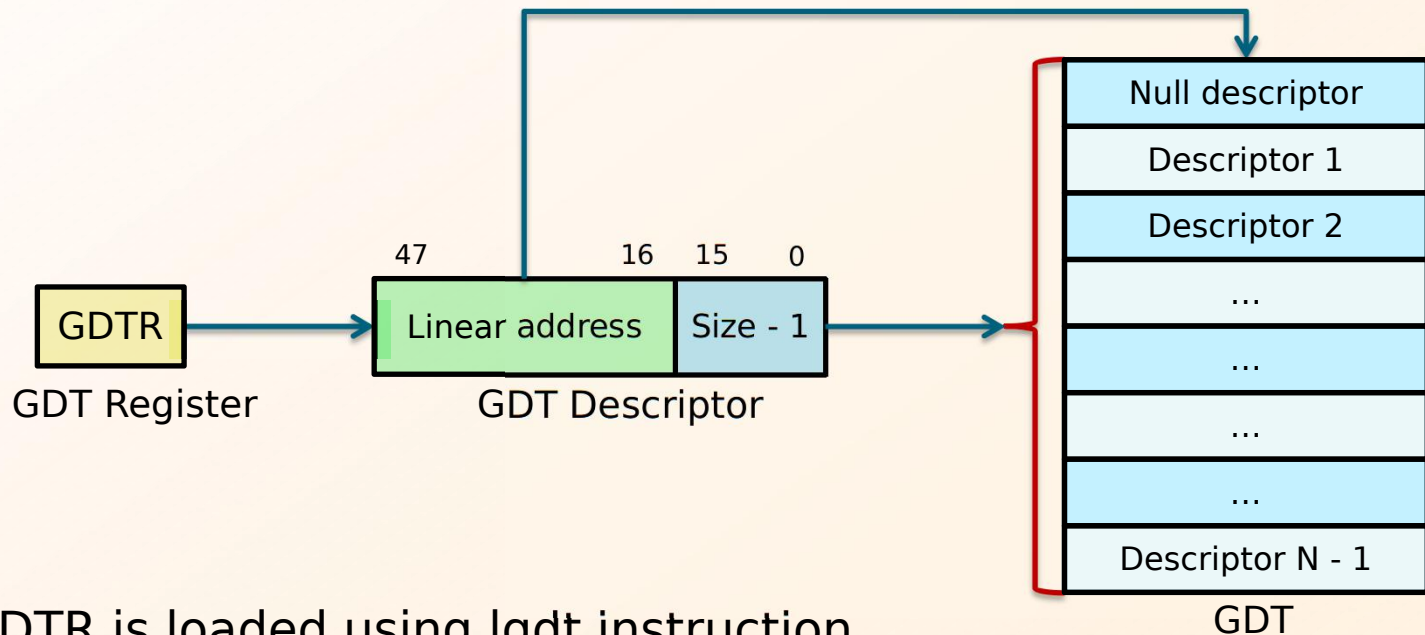
Global Descriptor Table (4)

- Access Byte
 - Present bit (P).
 - Descriptor Privilege Level (DPL).
 - Executable bit (E).
 - Direction/Conforming bit (D/C).
 - For stack/data segment, 0 = segment grows upwards, 1 = grows downwards.
 - If segment is code, 0 = non-conforming segment, 1 = conforming.
 - Readable/Writable bit (R/W). For code segment, 1 = readable, for data, 1 = writable.
 - Accessed bit (A).



Global Descriptor Table (5)

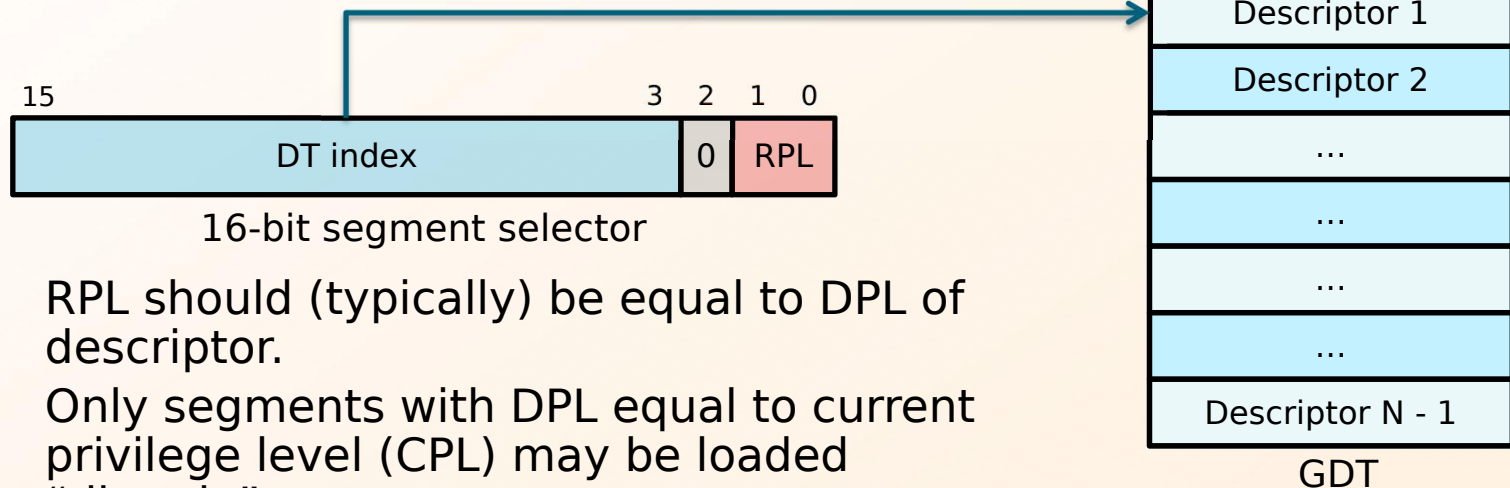
- The GDT is held in memory. The GDT register (GDTR) identifies where.



- GDTR is loaded using lgdt instruction.

What happens when loading segment register?

- It's complicated!
- So we ignore some of the hairy parts...
- A simplified overview ("half-truth"):



- RPL should (typically) be equal to DPL of descriptor.
- Only segments with DPL equal to current privilege level (CPL) may be loaded "directly".
- Base, limit and access rights of segment descriptor are stored in descriptor cache ("hidden" registers).
 - So GDT need not be accessed from memory repeatedly!


What about ring transitions?

- How do we change CPL if we only can transfer control directly to code segment with DPL equal to current CPL?
 - Trapping mechanisms
 - Interrupts (int/iret)
 - Exceptions
 - System calls (sysenter/sysleave)
 - Call gates (not used in modern OSs).


How to enable Protected Mode

1. Setup a valid GDT for Protected Mode.
2. Disable A20M to enable address line 20.
 - Allows us to use more than 1MB.
3. Disable interrupts.
4. Load GDT.
5. Enable Protection Enable (PE) bit (bit 0) in control register 0 (CR0).
6. Do a long jump to 32-bit code using valid code segment selector.
 - To clear CPU's prefetch input queue.
7. Setup valid data segment (DS) – and preferably ES, FS and GS too!
8. Setup valid stack segment (SS) and stack for Protected Mode.
9. That's it!

Any questions?



Any questions?



Resources

- <http://www.logix.cz/michal/doc/i386/chp06-03.htm>
- http://wiki.osdev.org/index.php?title=Global_Descriptor_Table&oldid=13668
- http://en.wikipedia.org/w/index.php?title=Global_Descriptor_Table&oldid=578074954
- <http://duartes.org/gustavo/blog/post/cpu-rings-privilege-and-protection>
- <http://www.osdever.net/tutorials/view/protected-mode>
- http://en.wikipedia.org/w/index.php?title=X86_memory_segmentation&oldid=581878645
- <http://www.rcollins.org/ddj/Aug98/Aug98.html>
- <http://www.win.tue.nl/~aeb/linux/kbd/A20.html>
- http://wiki.osdev.org/A20_Line